

Corpus-density gradient in CS-faculty NameRank. The firm, well-sampled contrast is USA > China > India (bold, $n \geq 10$, non-overlapping CIs); faded $n < 10$ rows are point estimates only. The gradient is corpus-density, not research-quality (Tsinghua/Peking rank low despite top-tier research reputations).

